



OpenMP – MPI Coordination




Issue



- Hybrid applications gaining popularity
 - Especially with many-core architectures
- Multiple programming libraries
 - Operate independently
 - No knowledge of presence of other libs or what those libs are doing
- Conflicts occur
 - Binding, resource utilization, ...
- OpenMP/MPI WG
 - Prototype use of PMIx for coordination



What is PMIx?




The Community



FUJITSU




<https://pmix.github.io/pmix>
<https://github.com/pmix>



v1.2 Features


- Typical startup operations
 - Put, get, commit, barrier, spawn, [dis]connect, publish/lookup
- System/job data returned upon PMIx_Init
 - Topology, proc locations, bindings, ...
 - Provided by host resource manager via PMIx
 - Full list: <https://github.com/PMIx/PMIx/wiki/2.8-Pmix-Server-Data-Requirements>



v2.0 Features

- Tool connections
 - Debugger, job submission, query
- Generalized query support
 - Job status, layout, system data, resource availability
- Logging
 - Status reports, error output
- Event notification
 - App, system generated
 - Subscribe, chained
 - Pre-emption, failures, timeout warning, ...
- Flexible allocations
 - Release resources, request resources
- Job control
 - Pause, kill, signal, heartbeat, resilience support


2Q2017




v3.0 Planned Features

- Network support
 - Security keys, pre-spawn local driver setup, fabric topology and status, traffic reports
- File system support
 - Dependency detection
 - Tiered storage caching strategies
- Inter-library coordination

4Q2017



PMIx Integration Architecture



Messenger not Doer

- Standardized APIs
 - Four he



- Convenience
 - Facilitates


- Serves as validation platform for

- Community

10¹⁸

Scheme

- Each library
 - Declares itself upon initialization
 - Model, implementation, version, ...
 - Registers event handler
 - Receive async notice of other declarations
 - Async updates as published by other libs
 - Order is irrelevant
 - Events cached and delivered upon registration
- Coordinate during execution via events
- Opportunity
 - Utilize PMIx features for OpenMP itself



Initialization

```
· pmix_info_t · *info;
· size_t · ninfo;
· pmix_status_t · code = PMIX_MODEL DECLARED;

· /* declare ourselves */
· ninfo = 4;
· PMIX_INFO_CREATE(info, · ninfo);
· PMIX_INFO_LOAD(&info[0], · PMIX_PROGRAMMING_MODEL, · "EXAMPLE", · PMIX_STRING);
· PMIX_INFO_LOAD(&info[1], · PMIX_MODEL_LIBRARY_NAME, · "FOOL", · PMIX_STRING);
· PMIX_INFO_LOAD(&info[2], · PMIX_MODEL_LIBRARY_VERSION, · "1.2.3", · PMIX_STRING);
· PMIX_INFO_LOAD(&info[3], · PMIX_THREADING_MODEL, · "NONE", · PMIX_STRING);
· if (PMIX_SUCCESS != (rc = PMIX_Init(&myproc, · info, · ninfo))) {
    . . . . .
    fprintf(stderr, · "PMIX_Init failed: · %s\n", · PMIX_Error_string(rc));
    . . . . .
    exit(1);
}
PMIX_INFO_FREE(info, · ninfo);

· /* register a handler specifically for when models declare */
· active = -1;
· ninfo = 1;
· PMIX_INFO_CREATE(info, · ninfo);
· PMIX_INFO_LOAD(&info[0], · PMIX_EVENT_HDLR_NAME, · "APP-MODEL", · PMIX_STRING);
· PMIX_Register_event_handler(&code, · 1, · info, · ninfo,
    . . . . .
    . . . . .
    model_callback, · model_registration_callback, · pointer_to_object);
· /* do whatever */
· . . . . .
· . . . . .
· . . . . .
· /* when registration completes */
· PMIX_INFO_FREE(info, · ninfo);
```

Callback when declarations received

Called back when
registration completes

object returned on
registration callback

Callback

- Receives
 - Copy of all info provided by declarer
 - `size_t` identifier of the event handler
 - Name of the event handler, if assigned
 - Provided object, if given
- Multiple handlers registered for same code(s)
 - Called in order (defined in registration)
 - Results of each handler aggregated and provided to later handlers

Summary

Working group:

Wiki: <https://github.com/pmix/pmix/wiki/7.1-openmp>

Minutes: <https://github.com/pmix/pmix/wiki/4.3-openmp>

Prototype code:

OpenMPI master: <https://github.com/open-mpi/ompi>

Example app:

<https://github.com/open-mpi/ompi/blob/master/orte/test/mpi/xlib.c>

We now have an interface library the RMs will support for application-directed requests

*Need to collaboratively define
what we want to do with it*

