

# HPC Resource Management: View to the Future

Ralph H. Castain

Jan 27, 2016



# Definition: RM

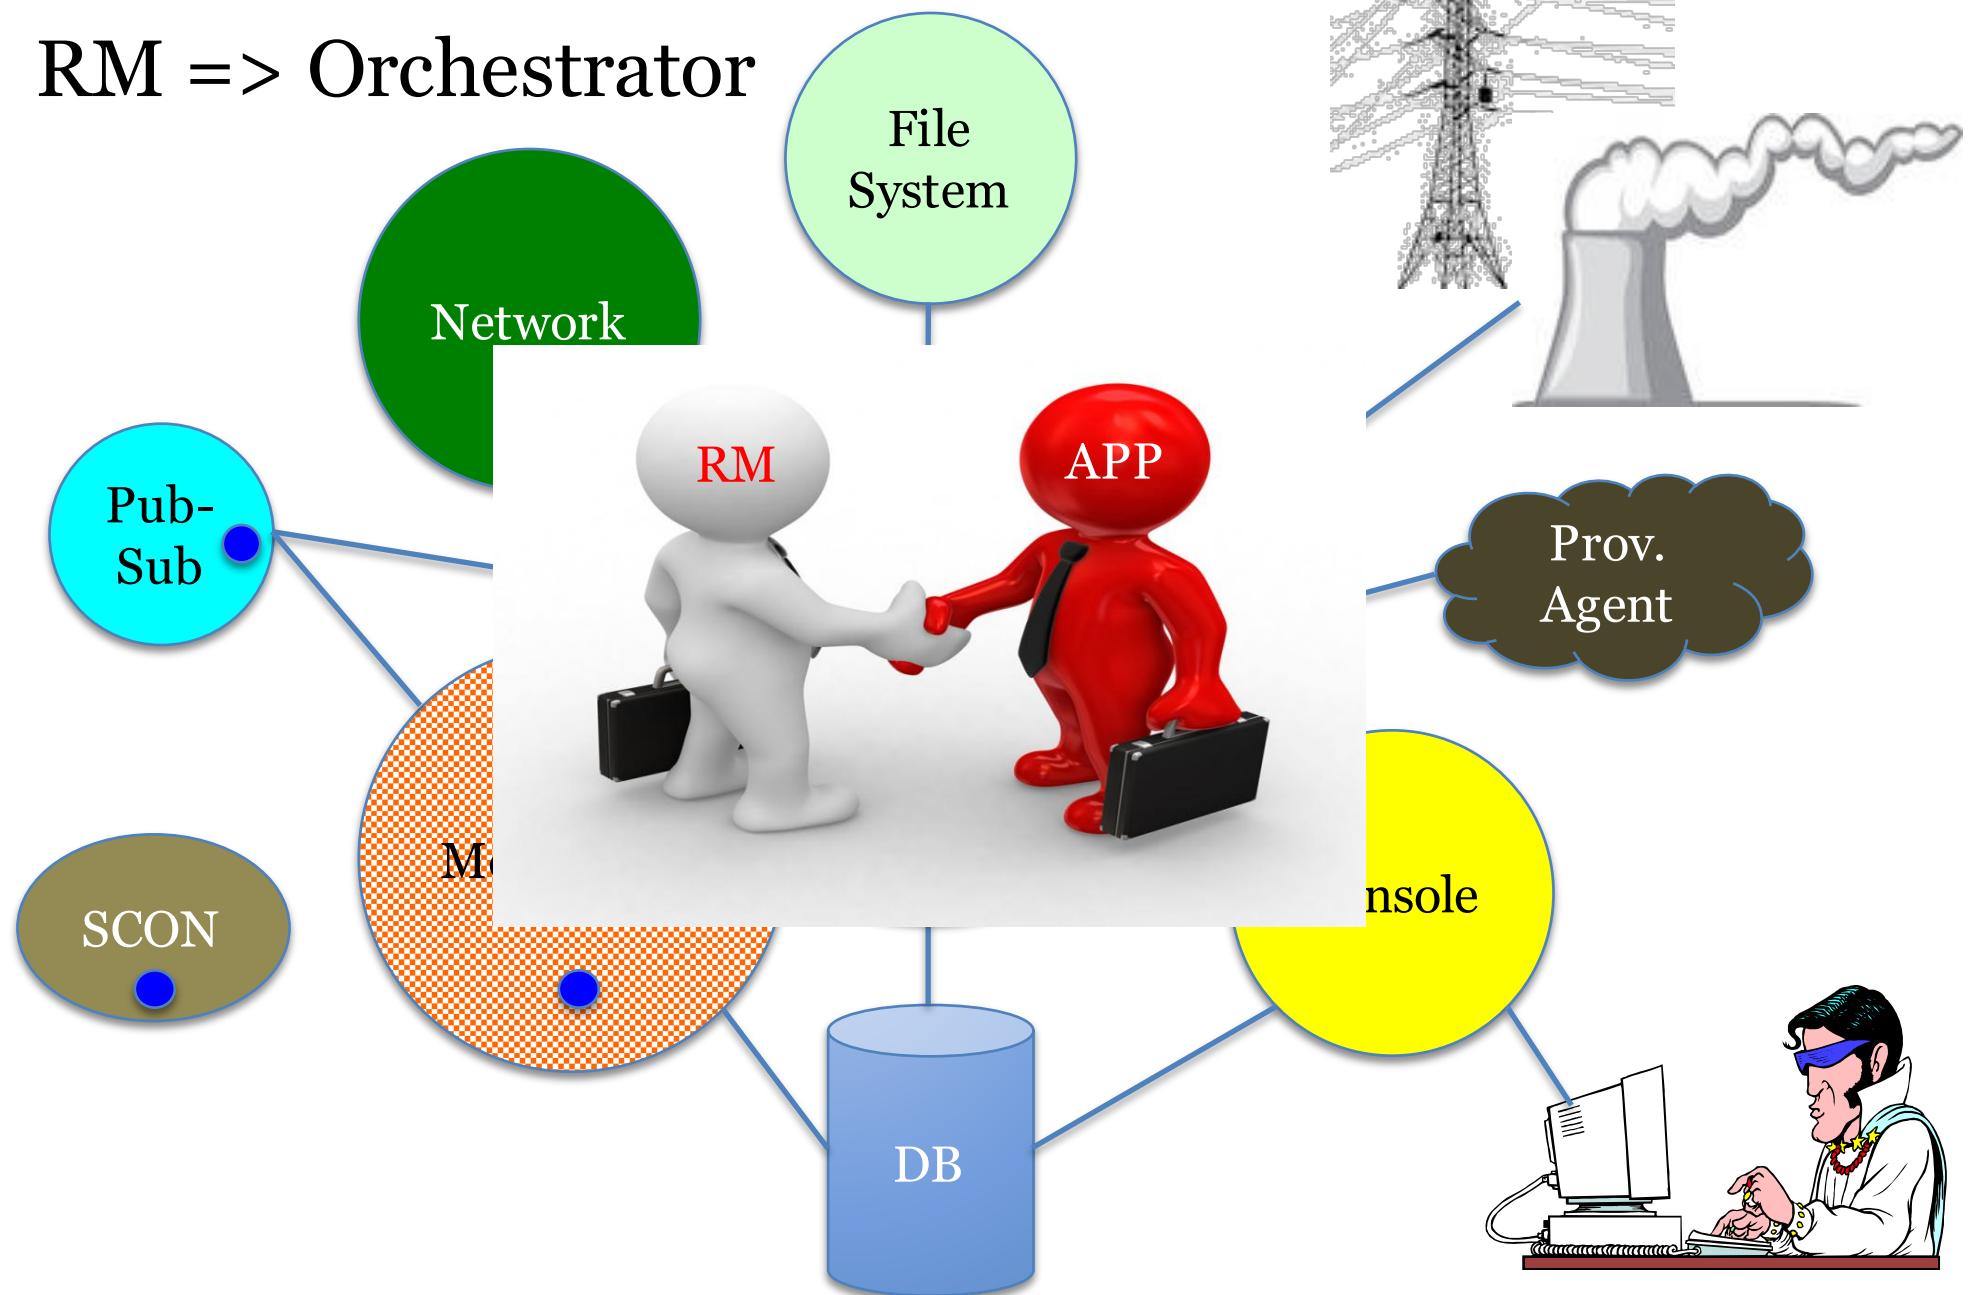
- Scheduler/Workload Manager (WLM)
  - Allocates resources to session
  - Interactive and batch
- Run-Time Environment
  - Launch and monitor applications
  - Support inter-process communication wireup
  - Serve as intermediary between applications and WLM
    - Dynamic resource requests
    - Error notification
  - Implement failure policies

# Current State

- Provided as complete package (WLM+RTE)
  - Sometimes offer hook to replace one
  - Mostly proprietary
  - Open source often GPL (integration difficult)
- Programming model specific, static
- Independent island
  - Not integrated with monitoring, file systems, or networks
- Limited fault tolerance support
  - Restart failed job



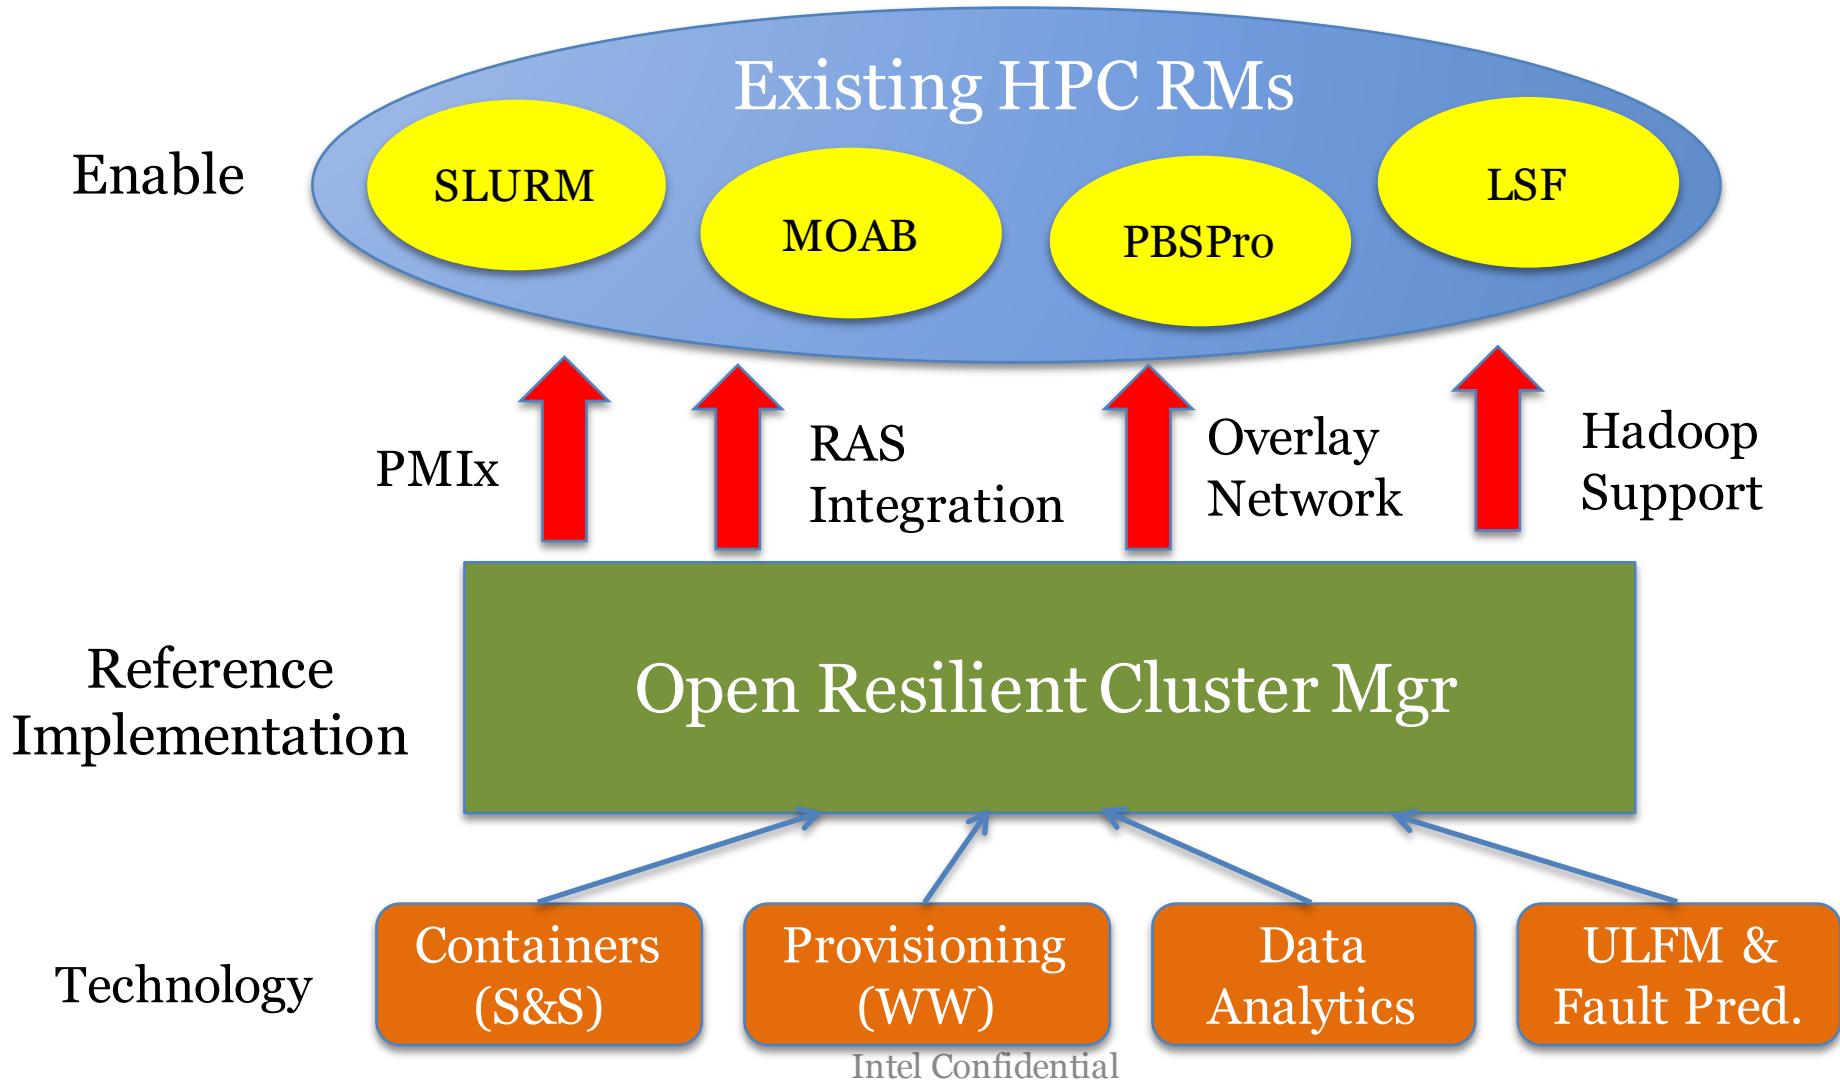
# RM => Orchestrator



# Resource Management Emerging Needs

- Emerging Application Needs
  - Dynamic resource management
  - Application involvement in decisions, including fault notification
  - Workflow orchestration
- Emerging System Needs
  - Scalable operations
  - Cross-subsystem integration (RM, RAS, site utilities, ...)
  - Data staging, Burst Buffers, persistent memory management
  - System notification requests, topology, QoS
  - Power management

# Multi-Tiered Strategy

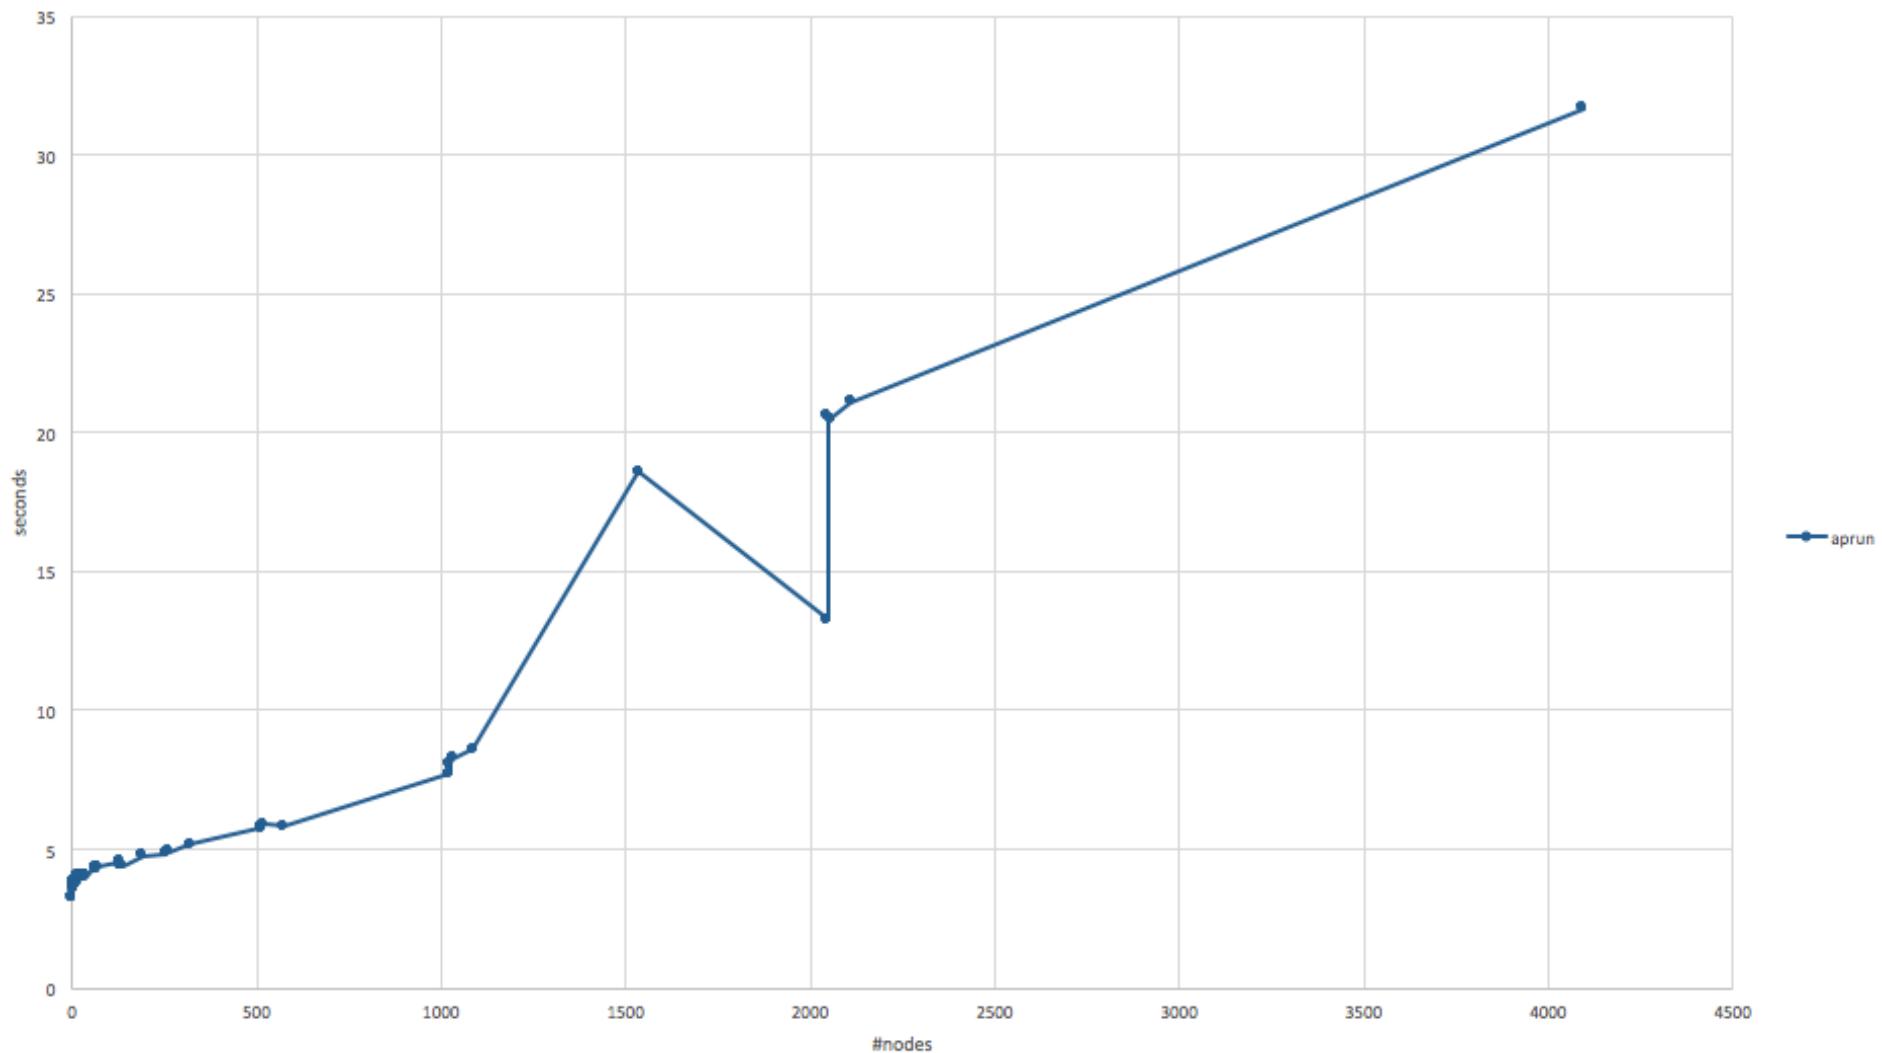


# Reference Implementation: ORCM

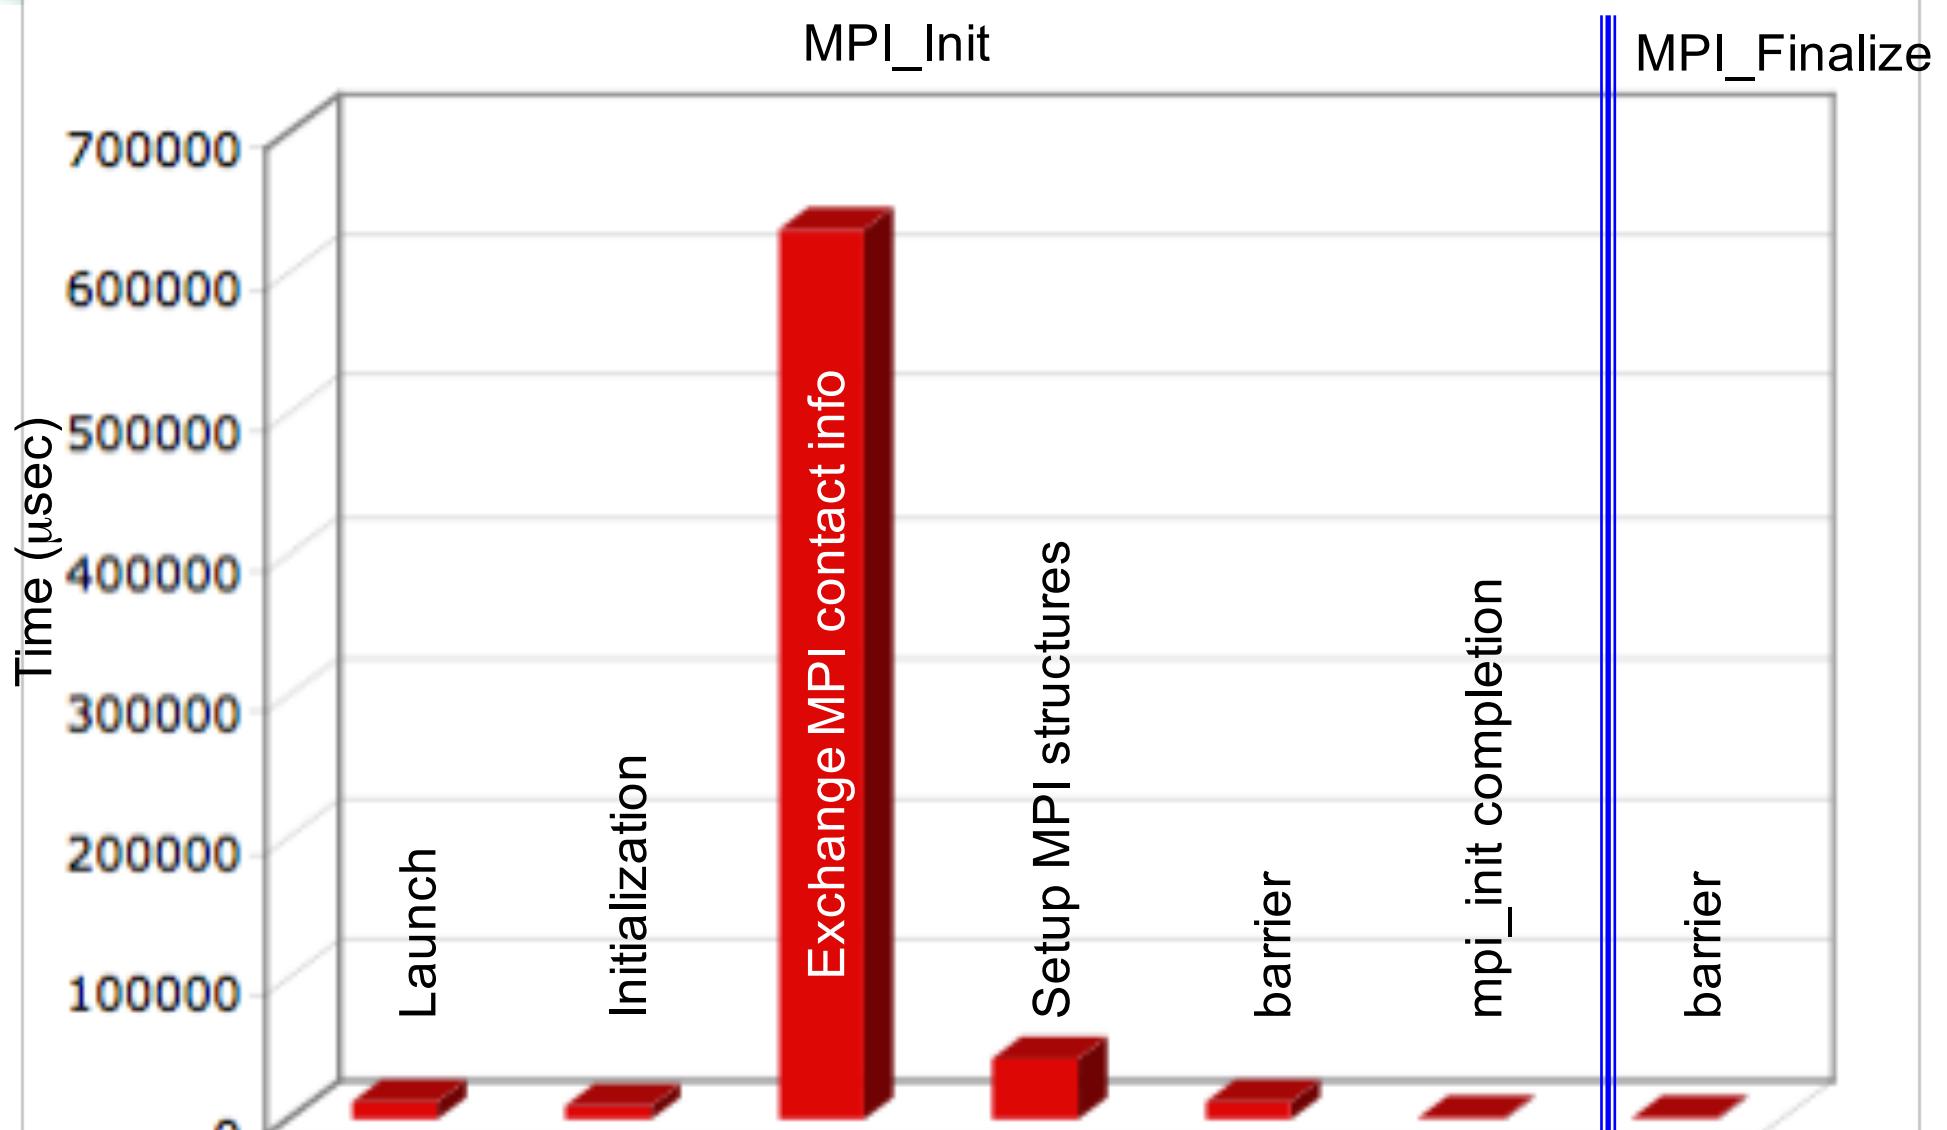
- Scalable to exascale levels & beyond
  - Better-than-linear scaling
  - Constrained memory footprint
- Dynamically configurable
  - Sense and adapt, user-directable
  - On-the-fly updates
  - Fully utilize underlying hw
- Open source (**non-GPL**)
  - Support proprietary add-ons
- Maintainable, flexible
  - Plugin architecture
  - Easy extension for R&D
- Resilient
  - Self-heal around failures
  - Reintegrate recovered resources

# Launch Scaling: Core Capability

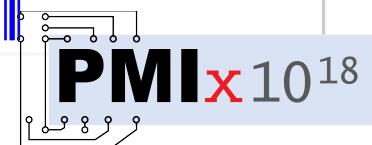
Time to launch and complete MPI\_Init/MPI\_Finalize



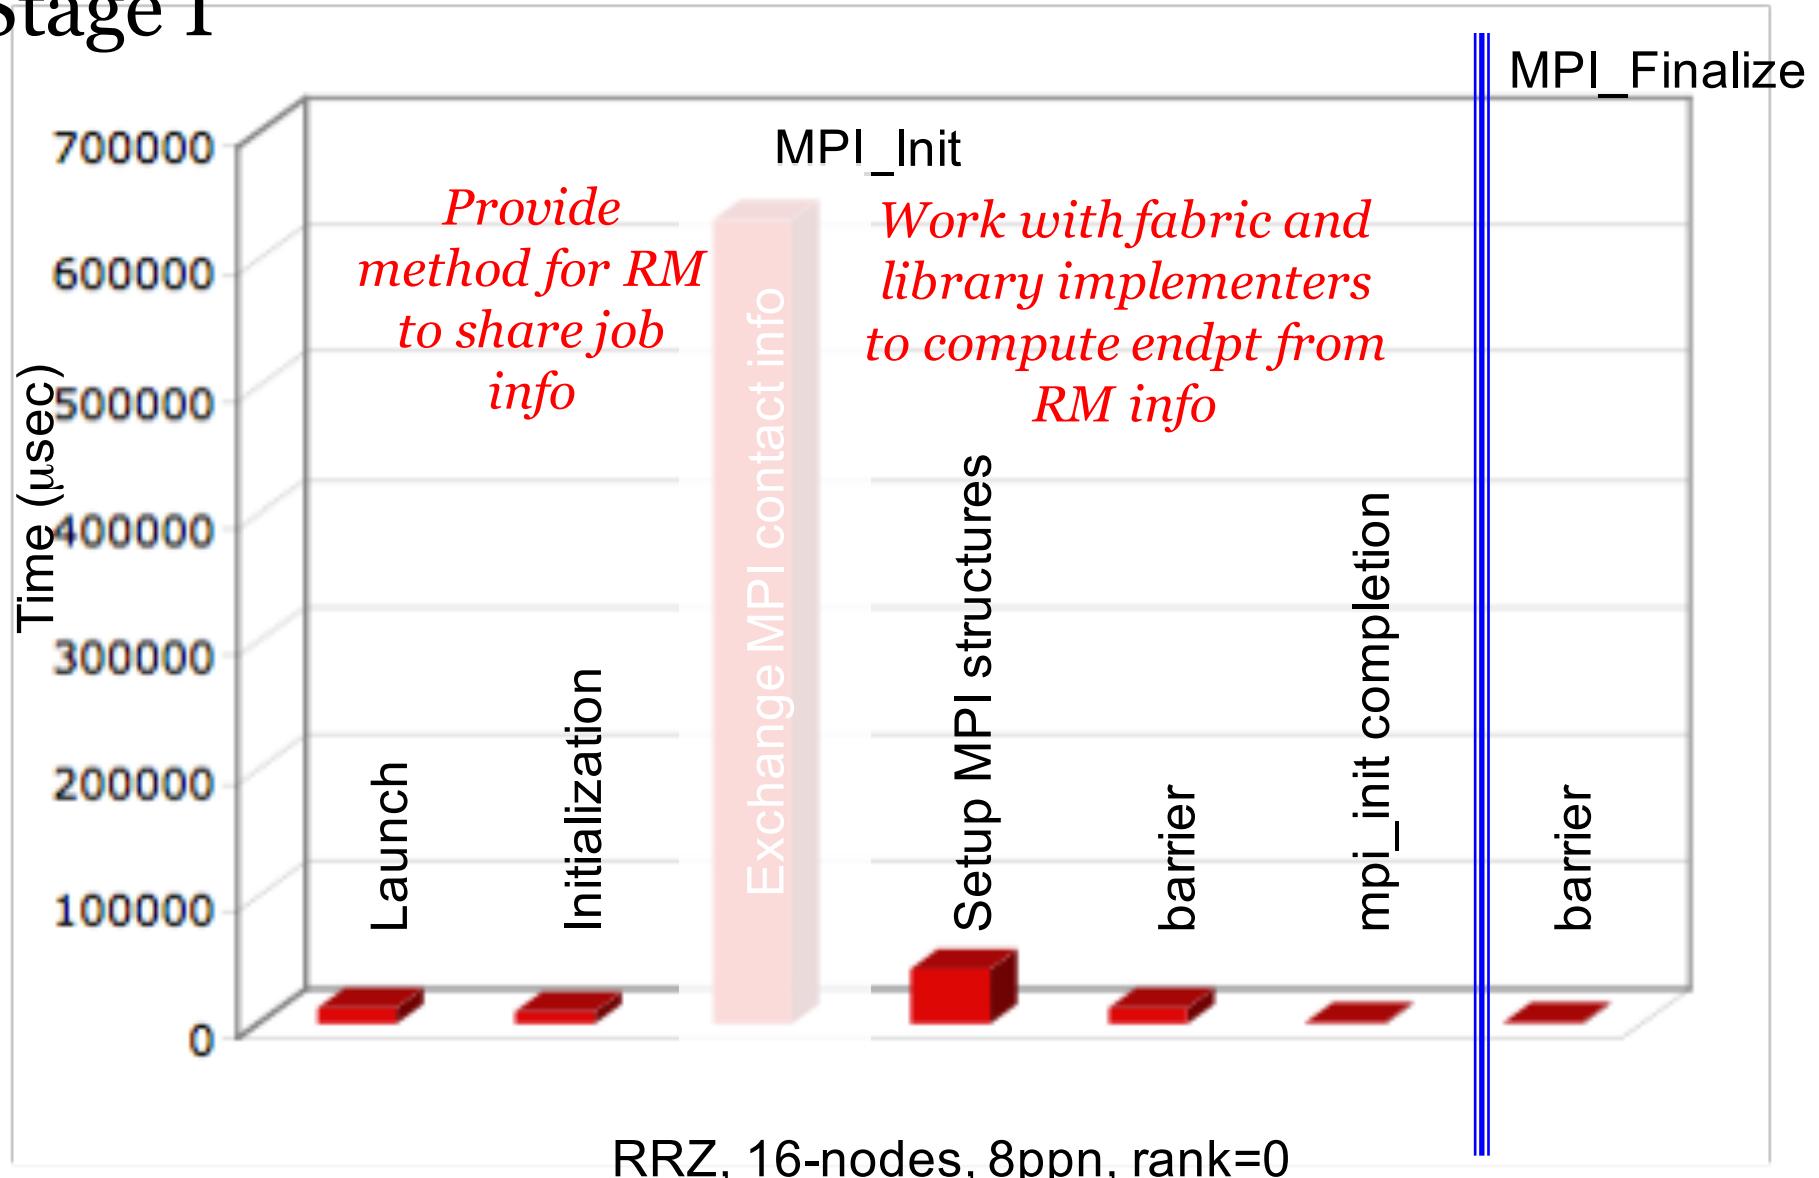
Intel Confidential



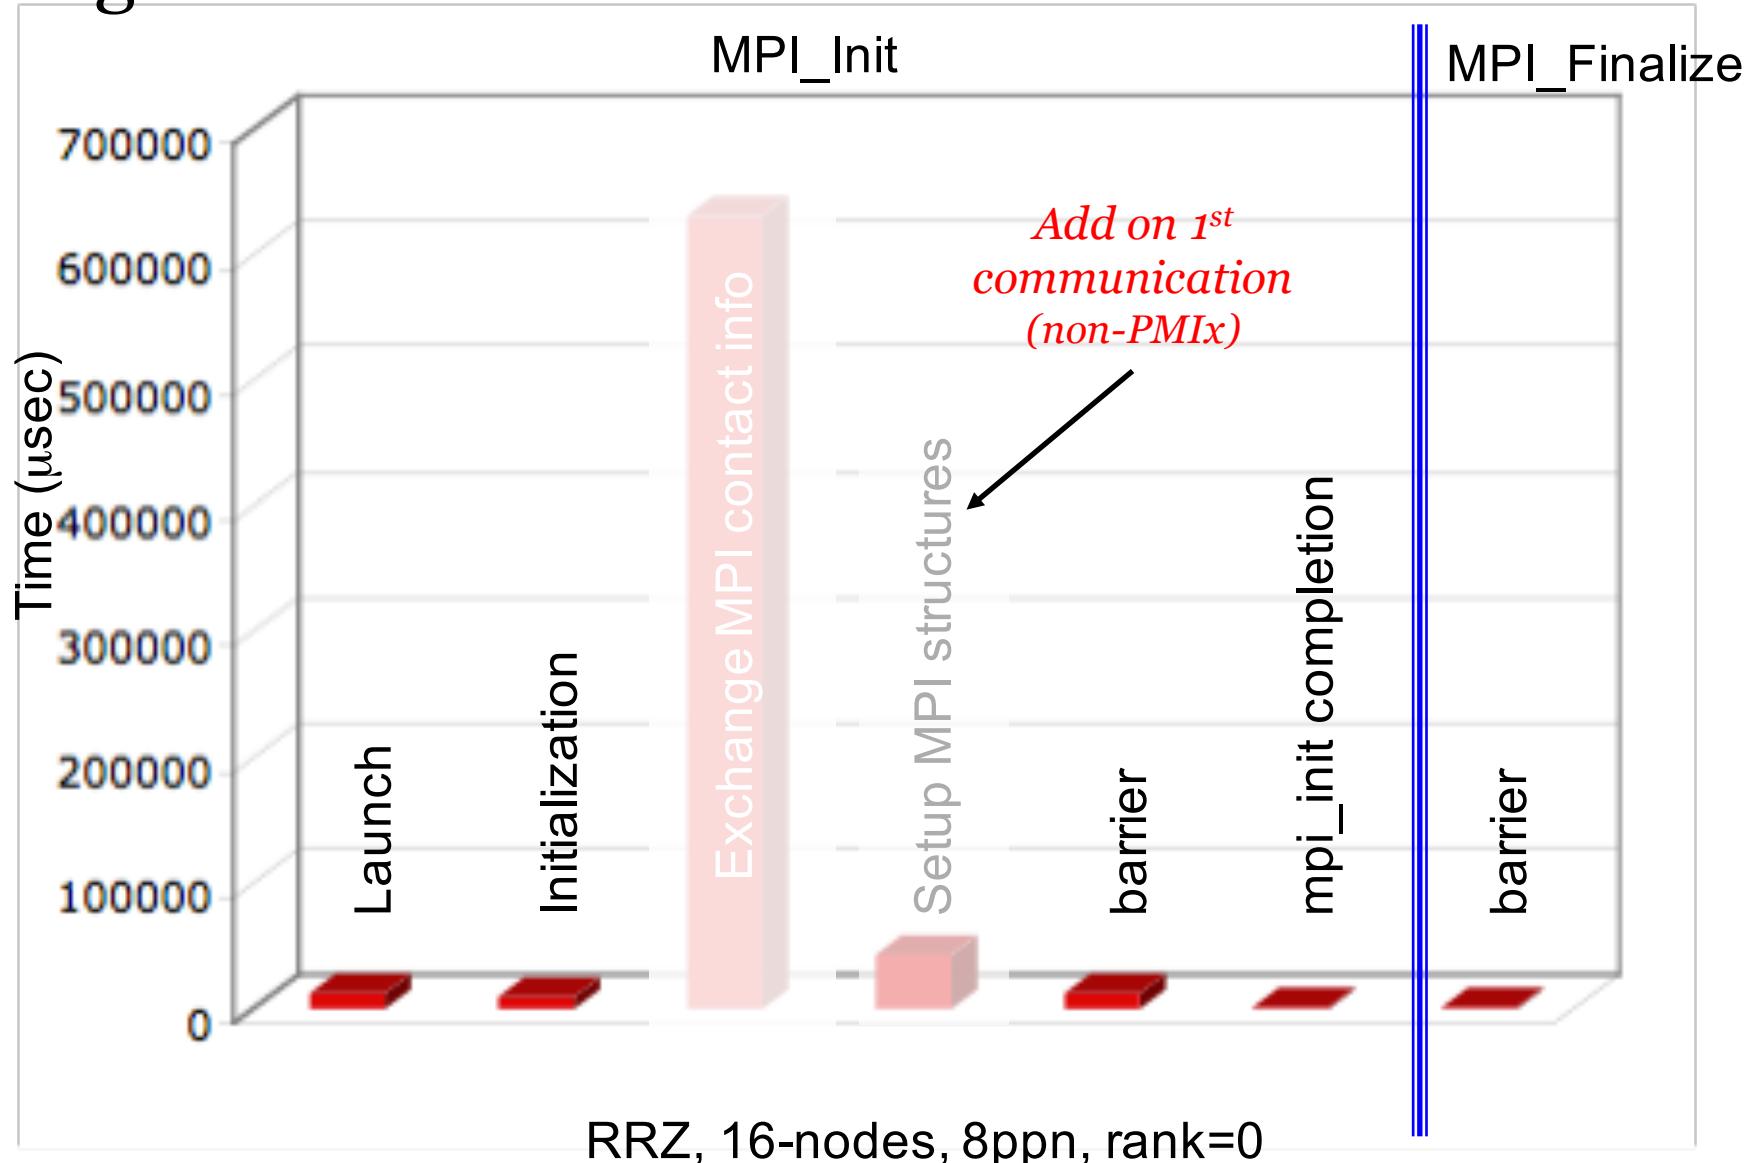
RRZ, 16-nodes, 8ppn, rank=0



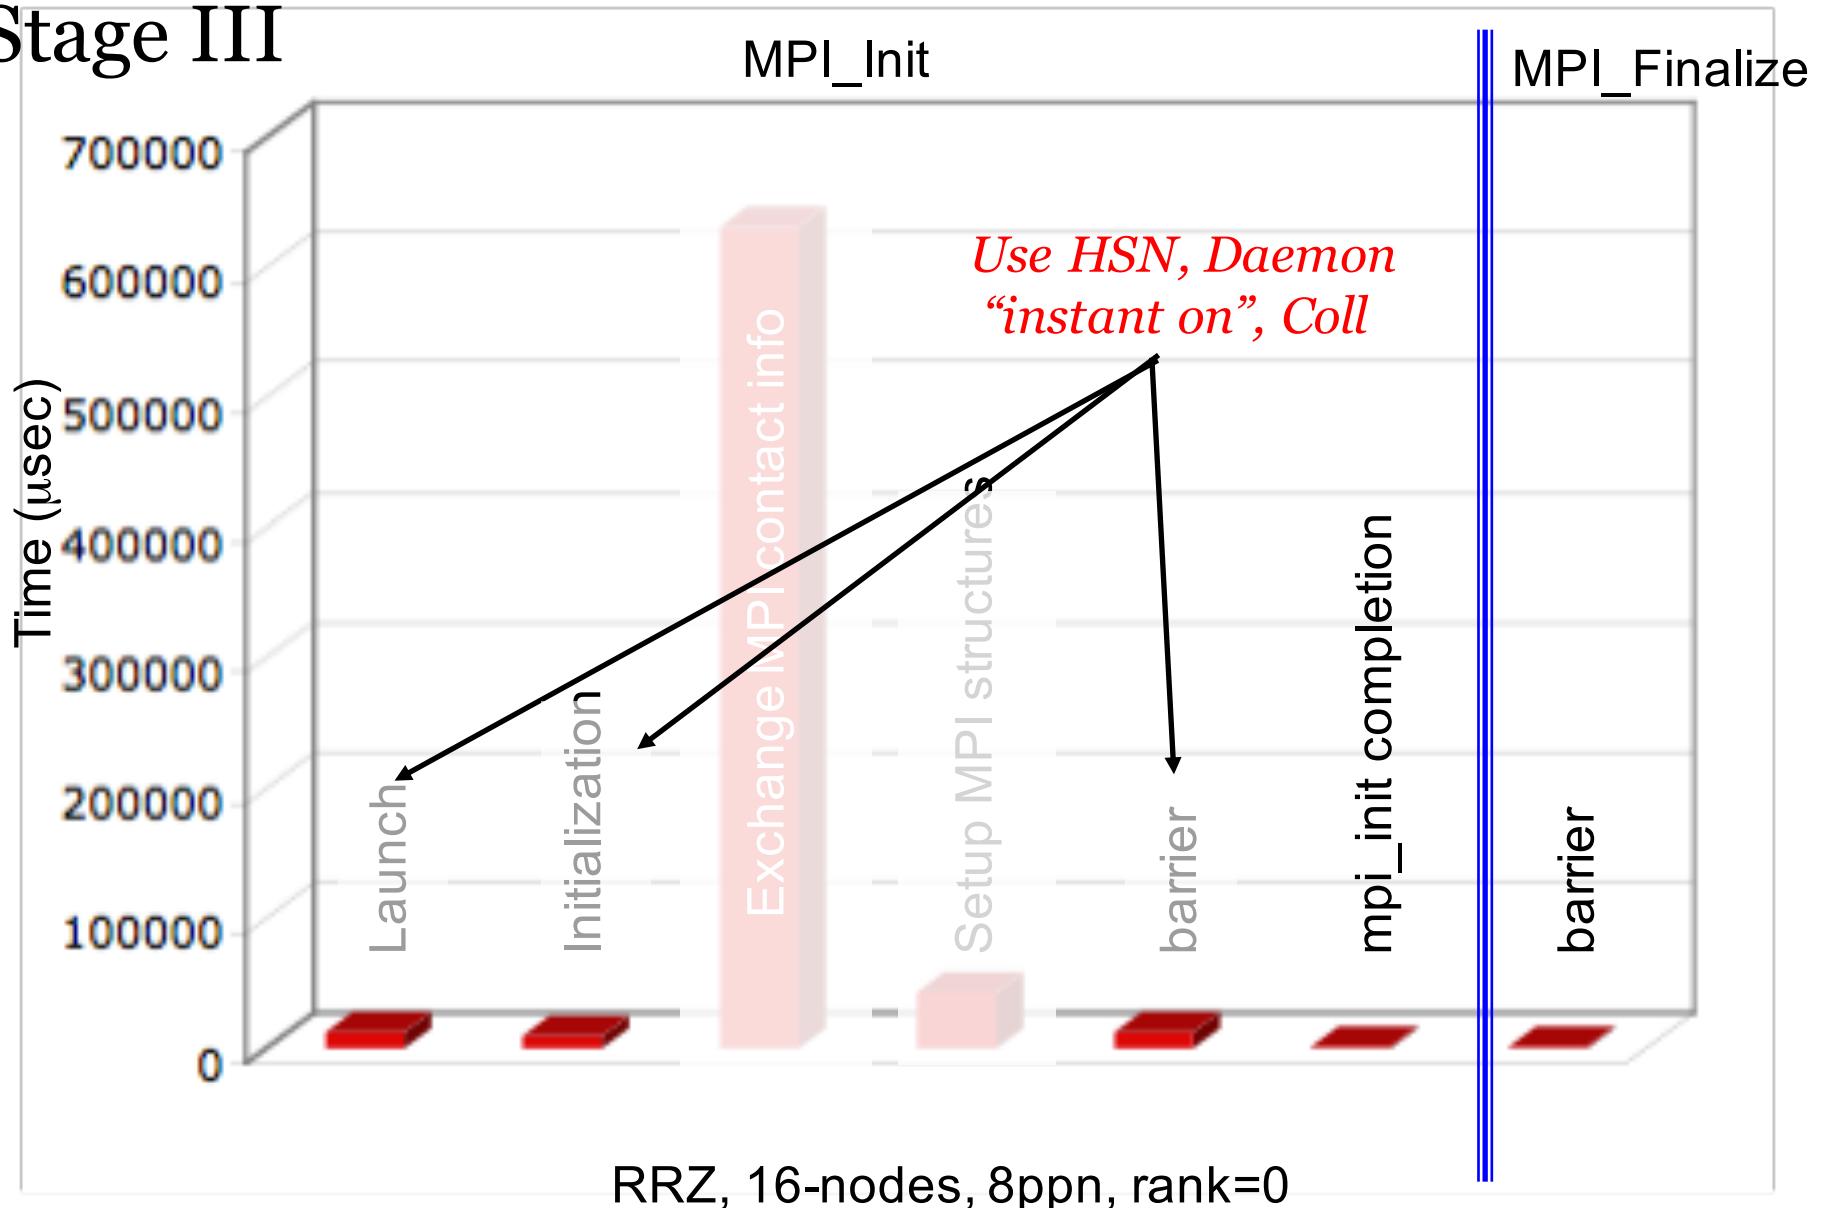
# Stage I



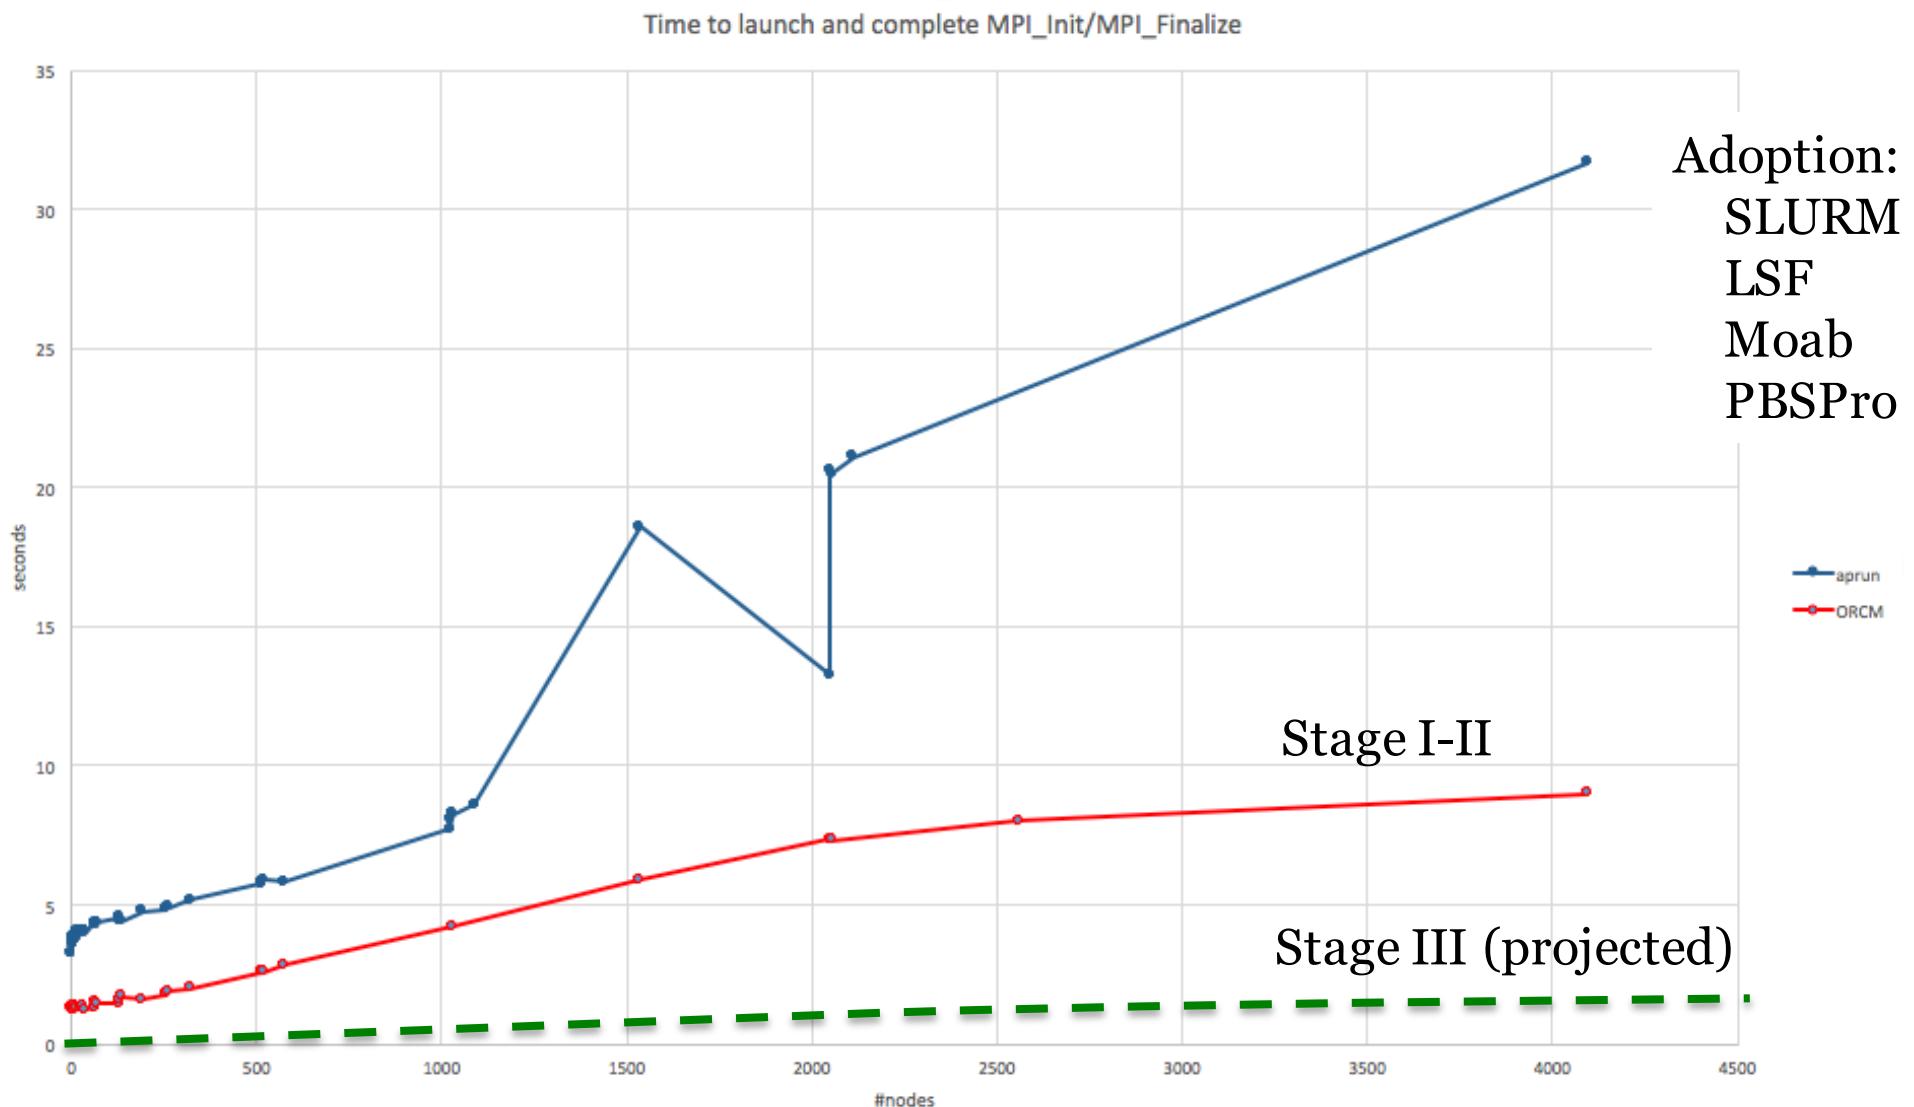
## Stage II



# Stage III



# Current Status



# Flexible Allocation Support

- Request additional resources
  - Compute, memory, network, NVM, burst buffer
  - Immediate, forecast
  - Expand existing allocation, separate allocation
- Return extra resources
  - No longer required
  - Will not be used for some specified time, reclaim (handshake) when ready to use
- Notification of preemption
  - Provide opportunity to cleanly pause

# I/O Support

- Asynchronous operations
  - Anticipatory data fetch, staging
  - Advise time to complete
  - Notify upon available
- Storage policy requests
  - Hot/warm/cold data movement
  - Desired locations and striping/replication patterns
  - Persistence of files, shared memory regions across jobs, sessions
  - ACL to generated data across jobs, sessions

# Spawn Support

- Staging support
  - Files, libraries required by new apps
  - Allow RM to consider in scheduler
    - Current location of data
    - Time to retrieve and position
    - Schedule scalable preload
    - Specified topology, min/max procs, ...
- Provisioning requests
  - Allow RM to consider in selecting resources, minimize startup time due to provisioning
  - Desired image, packages

# Network Integration

- Quality of service requests
  - Bandwidth, traffic priority, power constraints
  - Multi-fabric failover, striping prioritization
  - Security requirements
    - Network domain definitions, ACLs
- Notification requests
  - State-of-health
  - Update process endpoint upon fault recovery
- Topology information
  - Torus, dragonfly, ...

# Power Control/Management

- Scheduler analytics
  - Predict power usage, advise updates
- Application requests
  - Advise of changing workload requirements
  - Request changes in policy
  - Specify desired policy for spawned applications
- RM notifications
  - Need to change power policy
  - Preemption notification
    - Allow application to accept, request pause

# Fault Tolerance

- Notification
  - App can register for error notifications, incipient faults
    - RM-app negotiate to determine response
  - App can notify RM of errors
    - RM will notify specified, registered procs
- Rapid application-driven checkpoint
  - Local on-node NVRAM, auto-stripe checkpoints
  - Policy-driven “bleed” to remote burst buffers and/or global file system
- Restart support
  - Specify source (remote NVM checkpoint, global filesystem, etc)
  - Location hints/requests
  - Entire job, specific processes

# Workflow Orchestration

- Growing range of workflow patterns
  - Traditional HPC: bulk synchronous, single application
  - Analytics: asynchronous, staged
  - Parametric: large number of simultaneous independent jobs
- Ability to dynamically spawn a job
  - Need flexibility – min/max size, rolling allocation, ...
  - Events between jobs and cross-linking of I/O
  - Queuing of spawn requests
- Tool interfaces that work in these environments

# Workload Manager: Job Description Language

- Complexity of describing job is growing
  - Power, file/lib positioning
  - Performance vs capacity, programming model
  - System, project, application-level defaults
- Provide templates?
  - System defaults, with modifiers
    - --hadoop:mapper=foo,reducer=bar
  - User-defined
    - Application templates
    - Shared, group templates
  - Markup language definition of behaviors, priorities



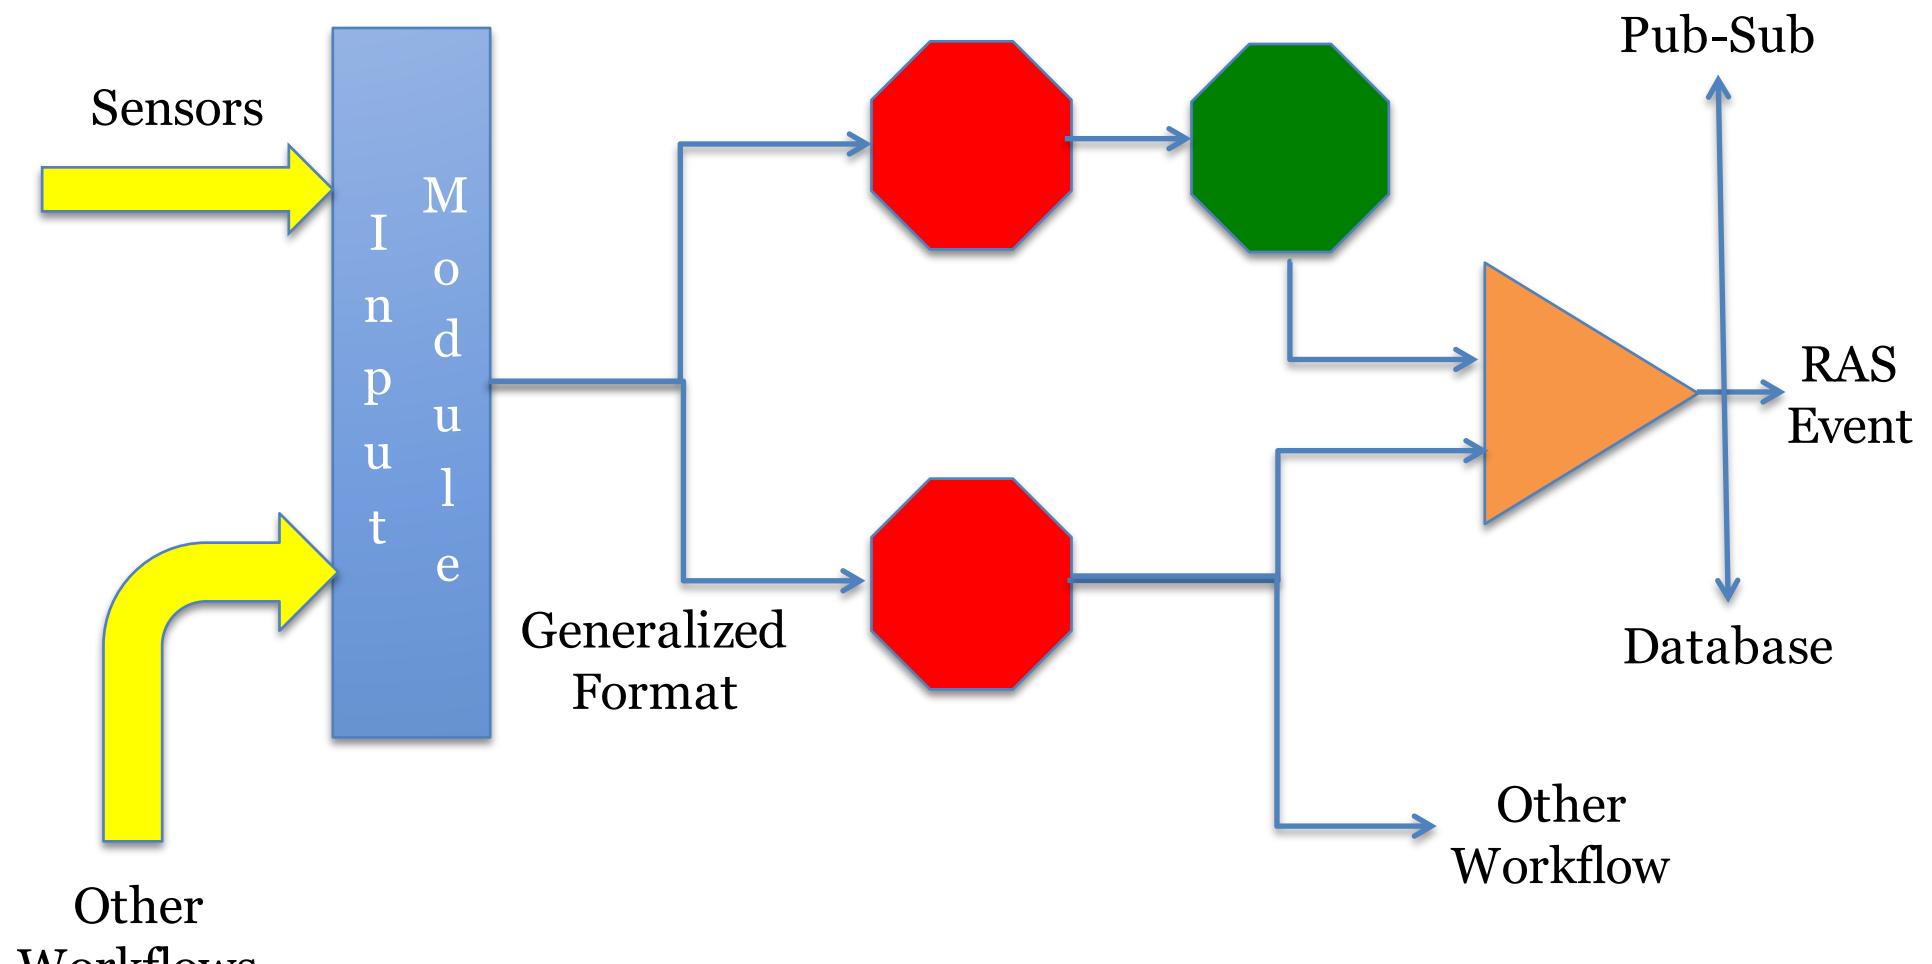
# Flexible Architecture

- Each tool built on top of same plugin system
  - Different combinations of frameworks
  - Different plugins activated to play different roles
  - Example: orcmd on compute node vs on rack/row controllers
- Designed for distributed, centralized, hybrid operations
  - Centralized for small clusters
  - Hybrid for larger clusters
  - Example: centralized scheduler, distributed “worker-bees”
- Accessible to users for interacting with RM
  - Add shim libraries (abstract, public APIs) to access framework APIs
  - Examples: SCON, pub-sub, in-flight analytics

# Breaking it Down

- Workload Manager
  - Dedicated framework
  - Plugins for two-way integration to external WM (Moab, Cobalt)
  - Plugins for implementing internal WM (FIFO)
- Run-Time Environment
  - Broken down into functional blocks, each with own framework
    - Loosely divided into three general categories
    - Messaging, launch, error handling
    - One or more frameworks for each category
  - Knitted together via “state machine”
    - Event-driven, async
    - Each functional block can be separate thread
    - Each plugin within each block can be separate thread(s)

# Analytics Workflow Concept



*Available in SCON as well*

# Workflow Elements

- Average (window, running, etc.)
- Rate (convert incoming data to events/sec)
- Threshold (high, low)
- Filter
  - Selects input values based on provided params
- RAS event
  - Generates a RAS event corresponding to input description
- Publish data

# Analytics

- Execute on aggregator nodes for in-flight reduction
  - Sys admin defines, user can define (if permitted)
- Event-based state machine
  - Each workflow in own thread, own instance of each plugin
  - Branch and merge of workflow
  - Tap stream between workflow steps
  - Tap data streams (sensors, others)
- Event generation
  - Generate events/alarms
  - Specify data to be included (window)

# Distributed Architecture

- Hierarchical, distributed approach for unlimited scalability
  - Utilize daemons on rack/row controllers
- Analysis done at each level of the hierarchy
  - Support rapid response to critical events
  - Distribute processing load
  - Minimize data movement
- RM's error manager framework controls response
  - Based on specified policies

# Fault Diagnosis

- Identify root cause and location
  - Sometimes obvious – e.g., when direct measurement
  - Other times non-obvious
    - Multiple cascading impacts
    - Cause identified by multi-sensor correlations (indirect measurement)
    - Direct measurement yields early report of non-root cause
    - Example: power supply fails due to borderline cooling + high load
- Estimate severity
  - Safety issue, long-term damage, imminent failure
- Requires in-depth understanding of hardware

# Fault Prediction: Methodology

- Exploit access to internals
  - Investigate optimal location, number of sensors
  - Embed intelligence, communications capability
- Integrate data from all available sources
  - Engineering design tests
  - Reliability life tests
  - Production qualification tests
- Utilize learning algorithms to improve performance
  - Both embedded, post process
  - Seed with expert knowledge

# Fault Prediction: Outcomes

- Continuous update of mean-time-to-preventative-maintenance
  - Feed into projected downtime planning
  - Incorporate into scheduling algo
- Alarm reports for imminent failures
  - Notify impacted sessions/applications
  - Plan/execute preemptive actions
- Store predictions
  - Algorithm improvement

# HPC Controls

